


Hi, my name is Anurag Dangi , I am currently working as DevOps Engineer at TCS with 2+ years of experience, where I manage and optimize CI/CD pipelines and application deployments using tools such as Jenkins, Docker, Nexus, SonarQube, and Kubernetes. I use GitHub for source code management and Maven for build automation. In total, I have 3 years of experience in the field.

Before joining TCS, I worked as an AWS Engineer - Trainee at Concentrix for 1 year, where I gained hands-on experience with AWS services including EC2, VPC, and S3.

Currently, I am part of a large team dedicated to enhancing our application to improve user satisfaction and operational efficiency. I hold a Bachelor's degree in Mechanical Engineering from Rajiv Gandhi Technological University, Bhopal, and I am presently pursuing a Master's degree in Cloud Computing from IIT Patna.


Devops: Devops is the union of people, process and technology to enable continues delivery of value to your end users.

DevOps is a way of working that combines software development and IT operations to improve teamwork, automate tasks, and deliver software faster and more reliably.



Total Cluster- 2
Total Master Node- 4
Total Worker Node- 36

Total Pods - 80+
Scaling Mechanism: **Horizontal Pod Autoscaler**
(HPA) based on CPU/memory usage.



Master Node Configuration- t3.large
Worker Node Configuration- t3.Medium

Containers per Pod: Typically 2-3 containers.

Project Name: Application Feature Enhancement

Objective: Our goal is to enhance the Application by refining existing features and adding new ones to boost both user satisfaction and operational efficiency.

Team: We have a robust team of 160 people working on this project. The Development Team, consisting of 25 members, focuses on coding and building the app, while the Operations Team, with 30 members, ensures everything runs smoothly, from quality assurance to deployment and troubleshooting.

Project Flow:

1. **Development Phase:** We start with Jira for managing tasks and issues. Developers use Eclipse to work on code, which is then committed to GitHub. This helps us keep everything organized and track progress.
2. **Continuous Integration (CI):**
 - **Building the Code:** We use Maven to compile and build the project.
 - **Code Quality:** SonarQube checks for any code issues or vulnerabilities.
 - **Security & Testing:** OWASP helps us spot security issues, while JUnit and Selenium ensure our code works as intended through automated testing.
3. **Continuous Deployment (CD):**
 - We deploy to various environments: DEV for initial testing, SIT for system integration, UAT for user feedback, and staging for final checks.
 - Kubernetes manages our containers, Docker packages our app, and Nexus keeps our Docker images organized.
4. **Production Deployment:** Deployments happen during off-peak hours to minimize user disruption. Each change is reviewed and approved before going live.
5. **Disaster Recovery (DR):**
 - We have regular backups and automated failover to handle any issues. Regular DR drills ensure our recovery plans are solid and up-to-date.

Monitoring & Post-Deployment:

- We use Prometheus and Grafana to monitor app performance and catch any issues early. Continuous checks help us address any problems quickly.

Infrastructure:

- The project runs on AWS Cloud, managed by IBM, with high-availability Kubernetes clusters ensuring reliable performance.

Project Overview:**Project Name: Application Feature Enhancement****Project Motive:**

The project focuses on enhancing the application to improve user experience and operational efficiency by updating existing functionalities and deploying new features.

Total Team: 160 people

- Application Team (40 people):
 - Team 1: Development (20 people)
 - 10 developers (including architects and a manager)
 - 5-6 DevOps engineers
 - Remaining team: Infrastructure specialists
 - Team 2: Operations (20 people)
 - QA Team: 3 people responsible for quality assurance testing
 - Testing Team: 2-3 people for system and integration testing
 - Monitoring Team: 2-3 people dedicated to continuous monitoring of applications and infrastructure
 - Deployment Team: 2-4 people in charge of pushing changes to different environments
 - Troubleshooting Team: 2-3 people handling incident and problem management

Pipeline Breakdown:

CI/CD Pipeline: Managed using Jenkins, with code versioning in GitHub and deployment orchestrated through Kubernetes and Docker.

1. Development Phase:

- Ticketing:
 - Jira is used for logging user requests, tasks, bugs, or feature requests.
 - Developers receive tickets based on priority set by management.
- Source Code Development:
 - Developers work on application code using Eclipse or VS Code.
 - Code changes are committed and pushed to GitHub.

2. Continuous Integration (CI) Stage:

- Build Stage (Maven):
 - Maven is used to compile and build the project.
- Static Code Analysis (SonarQube):
 - SonarQube scans the code for potential bugs, security vulnerabilities, and code smells.
- Security Testing (OWASP):
 - OWASP checks for common web application vulnerabilities such as SQL injection and cross-site scripting.
- Automated Testing (JUnit & Selenium):
 - JUnit performs unit testing on individual components.
 - Selenium performs end-to-end testing across different browsers and devices.

3. Continuous Deployment (CD) Stage:

- Environments in CD Pipeline:
 - DEV Environment: Initial deployment for developers to test changes.
 - SIT (System Integration Testing): For testing the integration of system components.
 - UAT (User Acceptance Testing): Final user acceptance tests to ensure business requirements are met.
 - Staging Environment: Replica of the production environment for final validation and performance testing.
- Tools Used in CD:
 - Kubernetes: Orchestrates containers across environments.
 - Docker: Containerizes the application for deployment.
 - Nexus: Manages Docker image artifacts.

4. Production Deployment:

- Non-working Hours Deployment (10 PM - 6 AM):
Deployments occur during non-working hours to minimize downtime and user impact.
- Approval Process:
 - Pull Request (PR) is raised, reviewed, and approved by managers before deployment by the deployment team.
- Production Environment:
 - Kubernetes cluster with 80 pods across 40 nodes (4 master nodes and 18 worker nodes per region).
- Pod Container Sizes:
 - Sized according to the specific needs of the microservice or application, optimized for resource utilization (CPU, memory).

5. Disaster Recovery (DR):

- DR Strategy:
 - Backup and Restore: Regular backups of critical data and configurations to ensure data restoration in case of loss.
 - Failover: Automated failover to a secondary region or data center to minimize downtime.
 - Testing: Regular DR drills and tests to ensure recovery procedures work as expected.
 - Documentation: Detailed DR plans maintained and updated to reflect any changes in infrastructure or applications.

Monitoring & Post-Deployment:

- Monitoring:
 - Prometheus and Grafana for monitoring and logging.
- Performance Testing:
 - Conducted in the Staging Environment to ensure performance under load.
- Continuous Monitoring:
 - Post-deployment monitoring checks performance, logging, and uptime. Issues are flagged and addressed by troubleshooting teams.

Infrastructure Details:

- Kubernetes Setup:
 - 2 high-availability clusters, each with 40 nodes and 80 pods.
 - Master Nodes: 4 master nodes distributed across regions.
 - Worker Nodes: 18 worker nodes per region.
- Cloud Infrastructure:
 - Managed on Oracle Cloud, handled by IBM.

Tools Overview:

- Jira: Ticketing and request management.
- GitHub: Source code version control.
- Jenkins: CI/CD pipeline management.
- Maven: Build automation.
- SonarQube: Static code analysis.
- OWASP: Security testing.
- JUnit: Unit testing.
- Selenium: Functional and UI testing.
- Docker: Containerization.
- Kubernetes: Container orchestration.
- Nexus: Artifact repository for Docker images.
- Prometheus & Grafana: Monitoring and logging.
- AWS Cloud: Cloud platform.
- IBM: Manages infrastructure.

AWS Ansible Kubernetes Terraform

questions->

<https://chatgpt.com/share/a90f8e35-0ce7-4449-9381-7259b3099dc3>

1. What I Have Done in Terraform

Answer: In my role as a DevOps Engineer, I have extensively used Terraform to automate the provisioning of infrastructure on AWS. I've written Terraform scripts to manage various AWS resources, including VPCs, EC2 instances, S3 buckets, and RDS databases. I've also leveraged Terraform modules to create reusable components, ensuring consistency and reducing the effort required for managing infrastructure across multiple environments.

2. AWS Services I've Used

Answer: I have hands-on experience with a wide range of AWS services, including:

- **Amazon VPC:** For creating isolated networks in the AWS cloud.
- **Amazon EC2:** For scalable computing power.
- **Amazon S3:** For object storage and backups.
- **AWS Lambda:** For serverless computing, running code in response to events.
- **Amazon RDS:** For managed relational databases.
- **AWS IAM:** For managing access to AWS resources.

3. VPC Creation

Answer: I have created multiple Amazon VPCs to provide isolated environments for different projects. This included setting up subnets across multiple availability zones for high availability, configuring route tables, and establishing secure communication using security groups and NACLs. The VPCs were designed to meet the client's specific requirements, such as private subnets for sensitive workloads and VPN connections for secure access from on-premises environments.

4. Security Group

Answer: I've managed AWS Security Groups to control inbound and outbound traffic to EC2 instances and other AWS resources. This involved writing Terraform configurations to define security group rules that allow or deny traffic based on IP ranges, protocols, and port numbers. Security Groups were essential in ensuring that only authorized traffic could reach the resources, providing a critical layer of security.

5. Requirements of Client for VPC Creation

Answer: Clients often require VPCs that are secure, scalable, and compliant with industry

regulations. Typical requirements include:

- **Network Segmentation:** Dividing the VPC into public and private subnets.
- **High Availability:** Distributing subnets across multiple availability zones.
- **Secure Access:** Implementing VPN or Direct Connect for secure communication between on-premises networks and AWS.
- **Access Control:** Using Security Groups and NACLs to enforce strict access policies.

6. Terraform Module

Answer: I have developed Terraform modules to standardize infrastructure deployment across projects. For example, I created a VPC module that includes configurations for subnets, route tables, security groups, and NAT gateways. This modular approach allows for easy replication of infrastructure in different environments, ensuring consistency and reducing the potential for errors.

7. Kubernetes Architecture

Answer: In Kubernetes, I've worked with an architecture that includes a control plane (with components like the API server, etcd, and scheduler) and worker nodes (running containerized applications). I've deployed and managed EKS (Elastic Kubernetes Service) clusters on AWS, ensuring proper configuration for high availability, autoscaling, and secure communication between services. This includes managing IAM roles for service accounts, setting up ingress controllers, and configuring persistent storage with EBS volumes.

8. Docker Volume

Answer: Docker volumes are crucial for persisting data in containerized environments. I've used Docker volumes to store database data, configuration files, and application logs outside the container's ephemeral storage. In AWS, I've integrated Docker volumes with EBS (Elastic Block Store) to provide persistent storage that remains available even when containers are stopped or replaced.

9. Difference Between RUN, CMD, ENTRYPOINT, and ENV in Docker

Answer:

- **RUN:** Executes commands at build time, creating a new layer in the Docker image. It's commonly used for installing packages or setting up the environment.
- **CMD:** Sets the default command to run when a container starts, which can be overridden by the user at runtime.
- **ENTRYPOINT:** Defines the main command to run when the container starts. Unlike CMD, ENTRYPOINT commands are not easily overridden at runtime.
- **ENV:** Sets environment variables inside the container, useful for configuring the application's runtime behavior.

10. What is Maven Workflow?

Answer: Maven is a build automation tool used primarily for Java projects. The typical workflow includes:

- **Compiling:** Converting source code into bytecode.
- **Testing:** Running automated tests to verify the code.
- **Packaging:** Creating a distributable JAR or WAR file.
- **Deploying:** Uploading the package to a repository or deploying it to an application server. In my projects, I've integrated Maven into CI/CD pipelines to automate these processes, ensuring that every commit triggers a build and test cycle, leading to faster and more reliable deployments.

11. AWS Kubernetes (EKS)

Answer: I have experience deploying and managing Kubernetes clusters using Amazon EKS. I've configured EKS clusters for high availability, utilizing multiple availability zones, and integrated with AWS services like ALB (Application Load Balancer) for ingress and IAM roles for managing permissions. I've also set up CI/CD pipelines that automatically deploy applications to EKS, ensuring a smooth transition from development to production.

12. Daily Duties in a Project

Answer: My daily duties typically involve:

- **Monitoring and managing CI/CD pipelines** to ensure smooth and efficient deployments.
- **Collaborating with development teams** to optimize deployment processes and infrastructure.
- **Automating infrastructure provisioning** using Terraform for AWS services.
- **Managing and monitoring Kubernetes clusters** to ensure applications are running optimally.
- **Responding to incidents** and trouble-shooting issues related to infrastructure, applications, and deployment processes.

13. Project Details

Answer: One of my key projects involved migrating a monolithic application to a microservices architecture using AWS. I was responsible for setting up the infrastructure on AWS using Terraform, deploying microservices in EKS, and implementing CI/CD pipelines using Jenkins. The project required close collaboration with the development team to containerize the application, establish secure networking, and integrate monitoring and logging solutions like CloudWatch and Prometheus.

14. AWS Serverless

Answer: I've used AWS serverless technologies, such as AWS Lambda, to build scalable, event-driven applications. For example, I developed a serverless data processing pipeline where AWS Lambda functions were triggered by S3 events to process and store data in DynamoDB. This approach allowed us to scale the application automatically based on the volume of incoming data without managing any servers.

15. CI/CD Pipeline

Answer: I have designed and implemented CI/CD pipelines using AWS services like CodePipeline, CodeBuild, and Jenkins. These pipelines automate the entire application lifecycle, from code commit to deployment. Key elements include:

- **Automated builds** using Maven or Gradle.
- **Running tests** to ensure code quality.
- **Deploying applications** to various environments, including staging and production.
- **Implementing rollback strategies** to recover quickly from failed deployments.
- **Integrating security and performance testing** into the pipeline to maintain high standards.

These answers would help you convey your AWS cloud experience and DevOps skills effectively during interviews.

Kubernetes installation- <https://chatgpt.com/-are/87f7a48c-8260-444b-9409-1978c3b1dd5d>
Sure, here are more detailed answers to each of the common Kubernetes (K8s) interview questions:

Interview question=> <https://chatgpt.com/-are/92bfc098-8e52-4e02-bf2e-e3151af5e50c>

1. Explain the Deployment Manifest File

A deployment manifest file in Kubernetes is a YAML or JSON file that defines the desired state of a deployment. It specifies how many replicas of an application should run, which container images to use, the container ports, environment variables, and other configurations. This file is used by Kubernetes to create and manage the deployment of applications. Here is an example of a basic YAML deployment manifest file:

yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-container
          image: my-image:latest
          ports:
            - containerPort: 80
```

2. Choreography vs Orchestration in Kubernetes

- **Choreography:** In a choreographed system, each microservice works independently and interacts with other services through events. Each service listens for events it is interested in and reacts accordingly. There is no central control; instead, the flow of communication is decentralized.
- **Orchestration:** In an orchestrated system, a central controller manages the interactions between services. This controller dictates how and when services should communicate with each other, providing a more structured and controlled approach.

3. Stateless vs Stateful Microservices

- **Stateless Microservices:** Do not retain any data or state between requests. Each request is independent, and any required state must be stored externally, such as in a database or a cache.
- **Stateful Microservices:** Retain state across multiple requests. They manage persistent data and require careful management of state, often using external storage systems to maintain consistency.

4. Role of Domain Events in Microservices

Domain events represent significant occurrences within a domain that can trigger reactions in other parts of the system or in other microservices. They help achieve decoupling between microservices, enabling them to react to changes and events in a more flexible and loosely-coupled manner.

5. Pod Communication in Kubernetes

Pods in Kubernetes communicate with each other primarily using the following mechanisms:

- **Service Discovery:** Kubernetes provides services with DNS names, making it easier for pods to find and communicate with each other.
- **Cluster IP:** Each pod gets an IP address, and services have a stable IP address within the cluster.
- **Environment Variables:** Kubernetes injects environment variables into pods to provide information about services.

6. Pod Security Policy

A Pod Security Policy (PSP) is a cluster-level resource in Kubernetes that controls security-sensitive aspects of the pod specification. PSPs are used to define and enforce security requirements, such as:

- Running as a non-root user
- Using specific security contexts
- Limiting access to host resources

7. Blue/Green Deployment Pattern

Blue/Green deployment is a technique to reduce downtime and risk by running two identical production environments, only one of which (blue) serves live production traffic. The other (green) is staged with the new version of the application. When the new version is ready and tested, traffic is switched to the green environment.

8. Securing Microservices

- **Authentication and Authorization:** Use protocols like OAuth2 and JWT.
- **Encryption:** Secure communication using TLS/SSL.
- **API Gateways:** Implement API gateways for security policies and rate limiting.
- **Service Mesh:** Use service meshes like Istio for managing and securing service-to-service communication.

9. Service in Kubernetes

A Service in Kubernetes is an abstraction that defines a logical set of pods and a policy to access them. There are three main types of services:

- **ClusterIP:** Exposes the service on an internal IP in the cluster.
- **NodePort:** Exposes the service on a static port on each node's IP.
- **LoadBalancer:** Exposes the service externally using a cloud provider's load balancer.

10. How Ingress Helps in Kubernetes

Ingress is a collection of rules that allow inbound connections to reach the cluster services. It provides:

- **Routing:** Directs traffic to the appropriate service based on HTTP routes.
- **SSL Termination:** Manages SSL/TLS certificates for encrypted communication.
- **Load Balancing:** Distributes traffic across multiple backend services.

11. API Versioning in Microservices

API versioning involves managing changes to APIs in a way that does not disrupt existing clients. It is crucial in microservices to ensure compatibility and smooth transitions when updates are made. Common versioning strategies include:

- **URL Path Versioning:** e.g., /v1/resource
- **Query Parameter Versioning:** e.g., /resource?version=1
- **Header Versioning:** Using custom headers to specify the API version.

12. Taints and Tolerations in Kubernetes

Taints and tolerations are mechanisms in Kubernetes that allow nodes to repel certain pods.

- **Taints:** Applied to nodes to mark them as having certain conditions.
- **Tolerations:** Applied to pods to allow them to schedule onto nodes with matching taints.

This helps ensure that certain pods only run on suitable nodes.

13. Handling Increased Traffic on Kubernetes Cluster

- **Horizontal Pod Autoscaling:** Automatically adjusts the number of pod replicas based on metrics like CPU usage.
- **Cluster Autoscaling:** Adds or removes nodes from the cluster based on resource utilization.
- **Load Balancing:** Distributes traffic efficiently across multiple pods and nodes.

14. Common Issues When Spinning Up a Container in Kubernetes

- **Image Pull Errors:** Issues with pulling the container image from the registry.
- **Resource Constraints:** Insufficient CPU or memory resources.
- **Networking Issues:** Problems with network policies, DNS resolution, or service discovery.

- **Configuration Errors:** Misconfigured deployment manifests or environment variables.

15. Comprehensive Backup Strategy for Kubernetes

- **Etdcd Backup:** Regularly backup the etcd datastore, which stores the cluster state.
- **Persistent Volume Backup:** Backup data stored in persistent volumes.
- **Configuration Backup:** Backup Kubernetes resources like deployments, services, and config maps using tools like kubectl or Velero.
- **Application-level Backup:** Use application-specific backup mechanisms for stateful applications.

Kubernetes Tutorial

1. Kubernetes Architecture

Describe the Kubernetes architecture:

- Kubernetes uses a master-worker architecture. The **control plane** components (API server, etcd, controller manager, scheduler) manage the cluster's state, while **worker nodes** run application workloads.
 - **API Server:** The front-end of the control plane that handles REST commands used to control the cluster.
 - **etcd:** A highly available key-value store for all cluster data.
 - **Controller Manager:** Manages routine tasks in the cluster, such as managing nodes and pods.
 - **Scheduler:** Assigns workloads to nodes based on resource availability.

Main components of the Kubernetes control plane:

- **API Server:** Exposes the Kubernetes API and serves as the entry point for managing the cluster.
- **etcd:** Stores all cluster data and maintains the cluster's state.
- **Controller Manager:** Runs background processes to ensure the desired state of the cluster.
- **Scheduler:** Determines which nodes will run the pods based on resource requirements and other constraints.

How nodes interact with the control plane:

- Nodes run the **Kubelet**, which ensures containers are running in a pod. They communicate with the API server for cluster management and use **kube-proxy** to manage network traffic between pods.

What is etcd and its role in Kubernetes?

- **etcd** is a distributed key-value store that maintains the state of the cluster, including configuration data, metadata, and the status of various cluster resources.

2. Kubernetes Resources

Write a YAML file for a Deployment:

yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
```

```

app: my-app
template:
  metadata:
    labels:
      app: my-app
  spec:
    containers:
      - name: my-container
        image: my-image:latest
        ports:
          - containerPort: 80

```

- **Purpose:** This YAML creates a Deployment named my-deployment with three replicas of a pod running my-image:latest.

Explain the differences between Deployments, ReplicaSets, and Pods:

- **Pods:** The smallest Kubernetes object representing a single instance of a running process.
- **ReplicaSets:** Ensures a specified number of pod replicas are running. It is used by Deployments to maintain the desired state.
- **Deployments:** Manages ReplicaSets and provides declarative updates for pods, allowing you to describe the desired state of an application and roll out updates.

Creating a Service in Kubernetes:

yaml

```

apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: my-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376

```

- **ClusterIP:** Exposes the service on a cluster-internal IP, accessible only within the cluster.
- **NodePort:** Exposes the service on each Node's IP at a static port, accessible from outside the cluster.
- **LoadBalancer:** Exposes the service externally using a cloud provider's load balancer.

Using Secrets and ConfigMaps in Kubernetes:

- **Secret Example:**

yaml

```

apiVersion: v1
kind: Secret
metadata:
  name: my-secret
type: Opaque
data:
  username: YWRtaW4=
  password: MWYyZDFlMmU2N2Rm

```

- **Purpose:** Stores sensitive information such as passwords.
- **ConfigMap Example:**

yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-config
data:
  key1: value1
  key2: value2
```

- **Purpose:** Stores non-sensitive configuration data in key-value pairs.

Explain Persistent Volumes (PV) and Persistent Volume Claims (PVC):

- **PersistentVolume Example:**

yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: /data
```

- **PersistentVolumeClaim Example:**

yaml

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
```

- **PV:** A piece of storage provisioned by an administrator or dynamically via Storage Classes.
- **PVC:** A user request for storage.

3. Rollback

How to roll back a Deployment:

- Use the command:

-

```
kubectl rollout undo deployment/<deployment-name>
```

- **Purpose:** Rolls back to the previous version of the Deployment.

Command to view the revision history of a Deployment:

- Use:

-

kubectl rollout history deployment/<deployment-name>

- **Purpose:** Shows the revision history of the Deployment.

Undo a specific Deployment revision:

- Use:

-

kubectl rollout undo deployment/<deployment-name> --to-revision=<revision-number>

- **Purpose:** Rolls back to a specific revision.

4. Troubleshooting

Debugging a pod in a CrashLoopBackOff state:

- Check logs:

-

kubectl logs <pod-name>

- Describe the pod:

-

kubectl describe pod <pod-name>

- Investigate application issues, resource limits, dependencies, or configuration problems.

Investigate a pod not starting:

- Check events:

-

kubectl describe pod <pod-name>

- Check logs:

-

kubectl logs <pod-name>

- Verify resource requests and limits, image accessibility, and volume mounts.

Handle network connectivity issues between pods:

- Check network policies and Service configurations.
- Verify DNS and network settings.
- Use kubectl exec to test connectivity inside the pods.

5. Common Errors

Common issues with Deployments and resolutions:

- **Image pull errors:** Verify image name, registry credentials, and network access.
- **Pods stuck in Pending:** Check node resources, PVC bindings, and node selectors.
- **Failed scheduling:** Ensure nodes meet resource requirements and verify taints, tolerations, and selectors.

Common reasons for pods being in a Pending state:

- Insufficient node resources.
- Unbound PVCs.
- Unsatisfied node selectors or affinity rules.
- Reached cluster resource quotas.

Handling image pull errors:

- Verify image repository details.
- Check registry availability and authentication.

- Ensure network connectivity to the registry.

6. Node Size

Determining the appropriate node size:

- Consider resource requirements (CPU, memory, storage), scalability, and monitoring metrics.
- Consult best practices for the specific workloads.

Factors influencing node sizing:

- Number of pods and resource requests.
- Resource utilization patterns.
- High availability and network considerations.

Scaling a Kubernetes cluster:

- Manually adjust node count via cloud interfaces or Kubernetes tools.
- Use Cluster Autoscaler for automatic scaling based on usage.

7. RBAC (Role-Based Access Control)

What is RBAC and its importance:

- RBAC regulates access to resources based on user roles, ensuring security and appropriate access levels.

Creating a Role and RoleBinding:

- **Role Example:**

yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
```

- **RoleBinding Example:**

yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-pods
  namespace: default
subjects:
- kind: User
  name: "jane"
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

- **Role:** Grants permissions within a namespace.
- **RoleBinding:** Assigns the Role to a user.

Difference between Role and ClusterRole:

- **Role:** Namespace-specific.
- **ClusterRole:** Cluster-wide, applicable across all namespaces.

Debugging RBAC issues:

- Use:

-

kubectl auth can-i get pods --as=user@example.com

- Check RoleBindings and ClusterRoleBindings for correct configuration.

This guide covers key Kubernetes concepts and practical examples to help you prepare for interviews or manage your Kubernetes environment effectively. Let me know if you need more details on any topic!

Kubernetes Commands and Concepts Tutorial

Basic Concepts

1. Kubernetes Cluster

- A collection of nodes (machines) that run containerized applications managed by Kubernetes.

2. Node

- A single machine (physical or virtual) within a Kubernetes cluster that runs the containerized workloads.

3. Pod

- The smallest and simplest Kubernetes object. A pod can host one or more containers, including storage resources and a unique network IP.

4. Namespace

- A mechanism to partition cluster resources between multiple users or projects, effectively creating virtual clusters within a single physical cluster.

5. Deployment

- A Kubernetes controller that manages a set of identical pods, providing declarative updates and ensuring the desired number of replicas are running.

6. Service

- An abstraction that defines a logical set of pods and a policy for accessing them. Services facilitate network access to the pods.

Basic Commands and Their Purposes

1. kubectl version

- Displays the version information for both the Kubernetes client and server.
- **Example:**

-

kubectl version

2. kubectl cluster-info

- Provides information about the Kubernetes cluster.
- **Example:**

-

kubectl cluster-info

3. kubectl get nodes

- Lists all nodes in the cluster.
- **Example:**

-

kubectl get nodes

4. **kubectl get pods**

- Lists all pods in the current namespace.
- **Example:**

-

kubectl get pods

5. **kubectl get services**

- Lists all services in the current namespace.
- **Example:**

-

kubectl get services

6. **kubectl get namespaces**

- Lists all namespaces in the cluster.
- **Example:**

-

kubectl get namespaces

Creating and Managing Resources

1. **Creating a Namespace**

- **Command:**

-

kubectl create namespace mynamespace

- **Purpose:** Creates a new namespace called mynamespace.

2. **Setting the Current Context to a Namespace**

- **Command:**

-

kubectl config set-context --current --namespace=mynamespace

- **Purpose:** Sets the default namespace to mynamespace for subsequent commands.

3. **Creating a Deployment**

- **YAML Configuration (my-deployment.yaml):**

yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: mydeployment

namespace: mynamespace

spec:

replicas: 2

selector:

matchLabels:

app: myapp

template:

metadata:

labels:

app: myapp

spec:

containers:

- name: nginx

image: nginx

- **Command to Apply the Deployment:**

-

kubectl apply -f my-deployment.yaml

- **Purpose:** Creates or updates a deployment named mydeployment in the mynamespace namespace.

4. **Scaling a Deployment**

- **Command:**

-

kubectl scale deployment mydeployment --replicas=3

- **Purpose:** Scales the deployment mydeployment to 3 replicas.

5. **Checking the Rollout Status of a Deployment**

- **Command:**

-

kubectl rollout status deployment mydeployment

- **Purpose:** Checks the status of the rollout for the deployment mydeployment.

Inspecting and Debugging

1. **Describing a Pod**

- **Command:**

-

kubectl describe pod mypod

- **Purpose:** Displays detailed information about the pod mypod.

2. **Fetching Logs from a Pod**

- **Command:**

-

kubectl logs mypod

- **Purpose:** Retrieves and displays logs from the container in the pod mypod.

3. **Executing a Command in a Pod**

- **Command:**

-

kubectl exec mypod -- /bin/ba-

- **Purpose:** Executes a command within the container of the pod mypod, such as opening a ba- -ell.

4. **Listing All Events**

- **Command:**

-

kubectl get events

- **Purpose:** Lists all events in the current namespace, useful for debugging.

Advanced Commands

1. **-ow Resource Usage of Nodes**

- **Command:**

-

kubecttl top nodes

- **Purpose:** Displays the current CPU and memory usage of nodes in the cluster.

2. **-ow Resource Usage of Pods**

- **Command:**

-

kubecttl top pods

- **Purpose:** Displays the current CPU and memory usage of pods in the cluster.

3. **List All Pods Across All Namespaces**

- **Command:**

-

kubecttl get pods --all-namespaces

- **Purpose:** Lists all pods across all namespaces for a cluster-wide view.

Example Workflow

1. **Create a Namespace**

- **Command:**

-

kubecttl create namespace mynamespace

2. **Switch to the New Namespace**

- **Command:**

-

kubecttl config set-context --current --namespace=mynamespace

3. **Create a Deployment**

- **Command:**

-

kubecttl apply -f my-deployment.yaml

4. **Verify the Deployment**

- **Command:**

-

kubecttl get deployments

5. **Scale the Deployment**

- **Command:**

-

kubecttl scale deployment mydeployment --replicas=3

6. **Check the Rollout Status**

- **Command:**

-

kubecttl rollout status deployment mydeployment

7. **Inspect the Pods**

- **Command:**

-

```
kubectl get pods
```

This guide would help you understand and effectively use basic to advanced Kubernetes commands. Feel free to ask if you have any more questions or need further clarification!

AWS Ansible Kubernetes Terraform

1. What I Have Done in Terraform

Answer: In my role as a DevOps Engineer, I have extensively used Terraform to automate the provisioning of infrastructure on AWS. I've written Terraform scripts to manage various AWS resources, including VPCs, EC2 instances, S3 buckets, and RDS databases. I've also leveraged Terraform modules to create reusable components, ensuring consistency and reducing the effort required for managing infrastructure across multiple environments.

2. AWS Services I've Used

Answer: I have hands-on experience with a wide range of AWS services, including:

- **Amazon VPC:** For creating isolated networks in the AWS cloud.
- **Amazon EC2:** For scalable computing power.
- **Amazon S3:** For object storage and backups.
- **AWS Lambda:** For serverless computing, running code in response to events.
- **Amazon RDS:** For managed relational databases.
- **AWS IAM:** For managing access to AWS resources.

3. VPC Creation

Answer: I have created multiple Amazon VPCs to provide isolated environments for different projects. This included setting up subnets across multiple availability zones for high availability, configuring route tables, and establishing secure communication using security groups and NACLs. The VPCs were designed to meet the client's specific requirements, such as private subnets for sensitive workloads and VPN connections for secure access from on-premises environments.

4. Security Group

Answer: I've managed AWS Security Groups to control inbound and outbound traffic to EC2 instances and other AWS resources. This involved writing Terraform configurations to define security group rules that allow or deny traffic based on IP ranges, protocols, and port numbers. Security Groups were essential in ensuring that only authorized traffic could reach the resources, providing a critical layer of security.

5. Requirements of Client for VPC Creation

Answer: Clients often require VPCs that are secure, scalable, and compliant with industry regulations. Typical requirements include:

- **Network Segmentation:** Dividing the VPC into public and private subnets.
- **High Availability:** Distributing subnets across multiple availability zones.
- **Secure Access:** Implementing VPN or Direct Connect for secure communication between on-premises networks and AWS.
- **Access Control:** Using Security Groups and NACLs to enforce strict access policies.

6. Terraform Module

Answer: I have developed Terraform modules to standardize infrastructure deployment across projects. For example, I created a VPC module that includes configurations for subnets, route tables, security groups, and NAT gateways. This modular approach allows for easy replication of infrastructure in different environments, ensuring consistency and reducing the potential for errors.

7. Kubernetes Architecture

Answer: In Kubernetes, I've worked with an architecture that includes a control plane (with components like the API server, etcd, and scheduler) and worker nodes (running containerized applications). I've deployed and managed EKS (Elastic Kubernetes Service) clusters on AWS, ensuring proper configuration for high availability, autoscaling, and secure communication between services. This includes managing IAM roles for service accounts, setting up ingress controllers, and configuring persistent storage with EBS volumes.

8. Docker Volume

Answer: Docker volumes are crucial for persisting data in containerized environments. I've used Docker volumes to store database data, configuration files, and application logs outside the container's ephemeral storage. In AWS, I've integrated Docker volumes with EBS (Elastic Block Store) to provide persistent storage that remains available even when containers are stopped or replaced.

9. Difference Between RUN, CMD, ENTRYPOINT, and ENV in Docker

Answer:

- **RUN:** Executes commands at build time, creating a new layer in the Docker image. It's commonly used for installing packages or setting up the environment.
- **CMD:** Sets the default command to run when a container starts, which can be overridden by the user at runtime.
- **ENTRYPOINT:** Defines the main command to run when the container starts. Unlike CMD, ENTRYPOINT commands are not easily overridden at runtime.
- **ENV:** Sets environment variables inside the container, useful for configuring the application's runtime behavior.

10. What is Maven Workflow?

Answer: Maven is a build automation tool used primarily for Java projects. The typical workflow includes:

- **Compiling:** Converting source code into bytecode.
- **Testing:** Running automated tests to verify the code.
- **Packaging:** Creating a distributable JAR or WAR file.
- **Deploying:** Uploading the package to a repository or deploying it to an application server. In my projects, I've integrated Maven into CI/CD pipelines to automate these processes, ensuring that every commit triggers a build and test cycle, leading to faster and more reliable deployments.

11. AWS Kubernetes (EKS)

Answer: I have experience deploying and managing Kubernetes clusters using Amazon EKS. I've configured EKS clusters for high availability, utilizing multiple availability zones, and integrated with AWS services like ALB (Application Load Balancer) for ingress and IAM roles for managing permissions. I've also set up CI/CD pipelines that automatically deploy applications to EKS, ensuring a smooth transition from development to production.

12. Daily Duties in a Project

Answer: My daily duties typically involve:

- **Monitoring and managing CI/CD pipelines** to ensure smooth and efficient deployments.
- **Collaborating with development teams** to optimize deployment processes and infrastructure.
- **Automating infrastructure provisioning** using Terraform for AWS services.
- **Managing and monitoring Kubernetes clusters** to ensure applications are running optimally.
- **Responding to incidents** and troubleshooting issues related to infrastructure, applications, and deployment processes.

13. Project Details

Answer: One of my key projects involved migrating a monolithic application to a microservices architecture using AWS. I was responsible for setting up the infrastructure on AWS using Terraform, deploying microservices in EKS, and implementing CI/CD pipelines using Jenkins. The project required close collaboration with the development team to containerize the application, establish secure networking, and integrate monitoring and logging solutions like CloudWatch and Prometheus.

14. AWS Serverless

Answer: I've used AWS serverless technologies, such as AWS Lambda, to build scalable, event-driven applications. For example, I developed a serverless data processing pipeline where AWS Lambda functions were triggered by S3 events to process and store data in DynamoDB. This approach allowed us to scale the application automatically based on the volume of incoming data without managing any servers.

15. CI/CD Pipeline

Answer: I have designed and implemented CI/CD pipelines using AWS services like CodePipeline, CodeBuild, and Jenkins. These pipelines automate the entire application lifecycle, from code commit to deployment. Key elements include:

- **Automated builds** using Maven or Gradle.
- **Running tests** to ensure code quality.
- **Deploying applications** to various environments, including staging and production.
- **Implementing rollback strategies** to recover quickly from failed deployments.
- **Integrating security and performance testing** into the pipeline to maintain high standards.

These answers could help you convey your AWS cloud experience and DevOps skills effectively during interviews.

Kubernetes installation- <https://chatgpt.com/-are/87f7a48c-8260-444b-9409-1978c3b1dd5d>
Sure, here are more detailed answers to each of the common Kubernetes (K8s) interview questions:

Interview question=> <https://chatgpt.com/-are/92bfc098-8e52-4e02-bf2e-e3151af5e50c>

1. Explain the Deployment Manifest File

A deployment manifest file in Kubernetes is a YAML or JSON file that defines the desired state of a deployment. It specifies how many replicas of an application should run, which container images to use, the container ports, environment variables, and other configurations. This file is used by Kubernetes to create and manage the deployment of applications. Here is an example of a basic YAML deployment manifest file:

yaml

Copy code

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: my-deployment
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    matchLabels:
```

```
      app: my-app
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
app: my-app
spec:
  containers:
  - name: my-container
    image: my-image:latest
    ports:
    - containerPort: 80
```

2. Choreography vs Orchestration in Kubernetes

- **Choreography:** In a choreographed system, each microservice works independently and interacts with other services through events. Each service listens for events it is interested in and reacts accordingly. There is no central control; instead, the flow of communication is decentralized.
- **Orchestration:** In an orchestrated system, a central controller manages the interactions between services. This controller dictates how and when services should communicate with each other, providing a more structured and controlled approach.

3. Stateless vs Stateful Microservices

- **Stateless Microservices:** Do not retain any data or state between requests. Each request is independent, and any required state must be stored externally, such as in a database or a cache.
- **Stateful Microservices:** Retain state across multiple requests. They manage persistent data and require careful management of state, often using external storage systems to maintain consistency.

4. Role of Domain Events in Microservices

Domain events represent significant occurrences within a domain that can trigger reactions in other parts of the system or in other microservices. They help achieve decoupling between microservices, enabling them to react to changes and events in a more flexible and loosely-coupled manner.

5. Pod Communication in Kubernetes

Pods in Kubernetes communicate with each other primarily using the following mechanisms:

- **Service Discovery:** Kubernetes provides services with DNS names, making it easier for pods to find and communicate with each other.
- **Cluster IP:** Each pod gets an IP address, and services have a stable IP address within the cluster.
- **Environment Variables:** Kubernetes injects environment variables into pods to provide information about services.

6. Pod Security Policy

A Pod Security Policy (PSP) is a cluster-level resource in Kubernetes that controls security-sensitive aspects of the pod specification. PSPs are used to define and enforce security requirements, such as:

- Running as a non-root user
- Using specific security contexts
- Limiting access to host resources

7. Blue/Green Deployment Pattern

Blue/Green deployment is a technique to reduce downtime and risk by running two identical production environments, only one of which (blue) serves live production traffic. The other (green) is staged with the new version of the application. When the new version is ready and tested, traffic is switched to the green environment.

8. Securing Microservices

- **Authentication and Authorization:** Use protocols like OAuth2 and JWT.
- **Encryption:** Secure communication using TLS/SSL.
- **API Gateways:** Implement API gateways for security policies and rate limiting.

- **Service Mesh:** Use service meshes like Istio for managing and securing service-to-service communication.

9. Service in Kubernetes

A Service in Kubernetes is an abstraction that defines a logical set of pods and a policy to access them. There are three main types of services:

- **ClusterIP:** Exposes the service on an internal IP in the cluster.
- **NodePort:** Exposes the service on a static port on each node's IP.
- **LoadBalancer:** Exposes the service externally using a cloud provider's load balancer.

10. How Ingress Helps in Kubernetes

Ingress is a collection of rules that allow inbound connections to reach the cluster services. It provides:

- **Routing:** Directs traffic to the appropriate service based on HTTP routes.
- **SSL Termination:** Manages SSL/TLS certificates for encrypted communication.
- **Load Balancing:** Distributes traffic across multiple backend services.

11. API Versioning in Microservices

API versioning involves managing changes to APIs in a way that does not disrupt existing clients. It is crucial in microservices to ensure compatibility and smooth transitions when updates are made. Common versioning strategies include:

- **URL Path Versioning:** e.g., /v1/resource
- **Query Parameter Versioning:** e.g., /resource?version=1
- **Header Versioning:** Using custom headers to specify the API version.

12. Taints and Tolerations in Kubernetes

Taints and tolerations are mechanisms in Kubernetes that allow nodes to repel certain pods.

- **Taints:** Applied to nodes to mark them as having certain conditions.
- **Tolerations:** Applied to pods to allow them to schedule onto nodes with matching taints.

This helps ensure that certain pods only run on suitable nodes.

13. Handling Increased Traffic on Kubernetes Cluster

- **Horizontal Pod Autoscaling:** Automatically adjusts the number of pod replicas based on metrics like CPU usage.
- **Cluster Autoscaling:** Adds or removes nodes from the cluster based on resource utilization.
- **Load Balancing:** Distributes traffic efficiently across multiple pods and nodes.

14. Common Issues When Spinning Up a Container in Kubernetes

- **Image Pull Errors:** Issues with pulling the container image from the registry.
- **Resource Constraints:** Insufficient CPU or memory resources.
- **Networking Issues:** Problems with network policies, DNS resolution, or service discovery.
- **Configuration Errors:** Misconfigured deployment manifests or environment variables.

15. Comprehensive Backup Strategy for Kubernetes

- **Etcd Backup:** Regularly backup the etcd datastore, which stores the cluster state.
- **Persistent Volume Backup:** Backup data stored in persistent volumes.
- **Configuration Backup:** Backup Kubernetes resources like deployments, services, and config maps using tools like kubectl or Velero.
- **Application-level Backup:** Use application-specific backup mechanisms for stateful applications.

Describe the Kubernetes architecture:

- Kubernetes uses a master-worker architecture. The **control plane** components (API server, etcd, controller manager, scheduler) manage the cluster's state, while **worker nodes** run application workloads.
 - **API Server:** The front-end of the control plane that handles REST commands used to control the cluster.
 - **etcd:** A highly available key-value store for all cluster data.
 - **Controller Manager:** Manages routine tasks in the cluster, such as managing nodes and pods.
 - **Scheduler:** Assigns workloads to nodes based on resource availability.

Main components of the Kubernetes control plane:

- **API Server:** Exposes the Kubernetes API and serves as the entry point for managing the cluster.
- **etcd:** Stores all cluster data and maintains the cluster's state.
- **Controller Manager:** Runs background processes to ensure the desired state of the cluster.
- **Scheduler:** Determines which nodes will run the pods based on resource requirements and other constraints.

How nodes interact with the control plane:

- Nodes run the **Kubelet**, which ensures containers are running in a pod. They communicate with the API server for cluster management and use **kube-proxy** to manage network traffic between pods.

What is etcd and its role in Kubernetes?

- **etcd** is a distributed key-value store that maintains the state of the cluster, including configuration data, metadata, and the status of various cluster resources.

2. Kubernetes Resources

Write a YAML file for a Deployment:

yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
      - name: my-container
        image: my-image:latest
        ports:
        - containerPort: 80
```

- **Purpose:** This YAML creates a Deployment named my-deployment with three replicas of a pod running my-image:latest.

Explain the differences between Deployments, ReplicaSets, and Pods:

- **Pods:** The smallest Kubernetes object representing a single instance of a running process.
- **ReplicaSets:** Ensures a specified number of pod replicas are running. It is used by Deployments to maintain the desired state.
- **Deployments:** Manages ReplicaSets and provides declarative updates for pods, allowing you to describe the desired state of an application and roll out updates.

Creating a Service in Kubernetes:

yaml

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: my-service
```

```
spec:
```

```
  selector:
```

```
    app: my-app
```

```
  ports:
```

```
    - protocol: TCP
```

```
      port: 80
```

```
      targetPort: 9376
```

- **ClusterIP:** Exposes the service on a cluster-internal IP, accessible only within the cluster.
- **NodePort:** Exposes the service on each Node's IP at a static port, accessible from outside the cluster.
- **LoadBalancer:** Exposes the service externally using a cloud provider's load balancer.

Using Secrets and ConfigMaps in Kubernetes:

- **Secret Example:**

yaml

```
apiVersion: v1
```

```
kind: Secret
```

```
metadata:
```

```
  name: my-secret
```

```
type: Opaque
```

```
data:
```

```
  username: YWRtaW4=
```

```
  password: MWYyZDFlMmU2N2Rm
```

- **Purpose:** Stores sensitive information such as passwords.
- **ConfigMap Example:**

yaml

```
apiVersion: v1
```

```
kind: ConfigMap
```

```
metadata:
```

```
  name: my-config
```

```
data:
```

```
  key1: value1
```

```
  key2: value2
```

- **Purpose:** Stores non-sensitive configuration data in key-value pairs.

Explain Persistent Volumes (PV) and Persistent Volume Claims (PVC):

- **PersistentVolume Example:**

yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: /data
```

- **PersistentVolumeClaim Example:**

yaml

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
```

- **PV:** A piece of storage provisioned by an administrator or dynamically via Storage Classes.
- **PVC:** A user request for storage.

3. Rollback

How to roll back a Deployment:

- Use the command:

-

kubectl rollout undo deployment/<deployment-name>

- **Purpose:** Rolls back to the previous version of the Deployment.

Command to view the revision history of a Deployment:

- Use:

-

kubectl rollout history deployment/<deployment-name>

- **Purpose:** Shows the revision history of the Deployment.

Undo a specific Deployment revision:

- Use:

-

kubectl rollout undo deployment/<deployment-name> --to-revision=<revision-number>

- **Purpose:** Rolls back to a specific revision.

4. Troubleshooting

Debugging a pod in a CrashLoopBackOff state:

- Check logs:

-

kubectl logs <pod-name>

- Describe the pod:

-

kubectl describe pod <pod-name>

- Investigate application issues, resource limits, dependencies, or configuration problems.

Investigate a pod not starting:

- Check events:

-

kubectl describe pod <pod-name>

- Check logs:

-

kubectl logs <pod-name>

- Verify resource requests and limits, image accessibility, and volume mounts.

Handle network connectivity issues between pods:

- Check network policies and Service configurations.
- Verify DNS and network settings.
- Use kubectl exec to test connectivity inside the pods.

5. Common Errors

Common issues with Deployments and resolutions:

- **Image pull errors:** Verify image name, registry credentials, and network access.
- **Pods stuck in Pending:** Check node resources, PVC bindings, and node selectors.
- **Failed scheduling:** Ensure nodes meet resource requirements and verify taints, tolerations, and selectors.

Common reasons for pods being in a Pending state:

- Insufficient node resources.
- Unbound PVCs.
- Unsatisfied node selectors or affinity rules.
- Reached cluster resource quotas.

Handling image pull errors:

- Verify image repository details.
- Check registry availability and authentication.
- Ensure network connectivity to the registry.

6. Node Size

Determining the appropriate node size:

- Consider resource requirements (CPU, memory, storage), scalability, and monitoring metrics.
- Consult best practices for the specific workloads.

Factors influencing node sizing:

- Number of pods and resource requests.
- Resource utilization patterns.
- High availability and network considerations.

Scaling a Kubernetes cluster:

- Manually adjust node count via cloud interfaces or Kubernetes tools.
- Use Cluster Autoscaler for automatic scaling based on usage.

7. RBAC (Role-Based Access Control)

What is RBAC and its importance:

- RBAC regulates access to resources based on user roles, ensuring security and appropriate access levels.

Creating a Role and RoleBinding:

- **Role Example:**

yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
```

- **RoleBinding Example:**

yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-pods
  namespace: default
subjects:
- kind: User
  name: "jane"
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
  • Role: Grants permissions within a namespace.
  • RoleBinding: Assigns the Role to a user.
```

Difference between Role and ClusterRole:

- **Role:** Namespace-specific.
- **ClusterRole:** Cluster-wide, applicable across all namespaces.

Debugging RBAC issues:

- Use:

-

kubectl auth can-i get pods --as=user@example.com

- Check RoleBindings and ClusterRoleBindings for correct configuration.

Kubernetes Commands and Concepts Tutorial

Basic Concepts

7. Kubernetes Cluster

- A collection of nodes (machines) that run containerized applications managed by Kubernetes.

8. **Node**

- A single machine (physical or virtual) within a Kubernetes cluster that runs the containerized workloads.

9. **Pod**

- The smallest and simplest Kubernetes object. A pod can host one or more containers, including storage resources and a unique network IP.

10. **Namespace**

- A mechanism to partition cluster resources between multiple users or projects, effectively creating virtual clusters within a single physical cluster.

11. **Deployment**

- A Kubernetes controller that manages a set of identical pods, providing declarative updates and ensuring the desired number of replicas are running.

12. **Service**

- An abstraction that defines a logical set of pods and a policy for accessing them. Services facilitate network access to the pods.

Basic Commands and Their Purposes

7. **kubectl version**

- Displays the version information for both the Kubernetes client and server.
- **Example:**

-

kubectl version

8. **kubectl cluster-info**

- Provides information about the Kubernetes cluster.
- **Example:**

-

kubectl cluster-info

9. **kubectl get nodes**

- Lists all nodes in the cluster.
- **Example:**

-

kubectl get nodes

10. **kubectl get pods**

- Lists all pods in the current namespace.
- **Example:**

-

kubectl get pods

11. **kubectl get services**

- Lists all services in the current namespace.
- **Example:**

-

kubectl get services

12. **kubectl get namespaces**

- Lists all namespaces in the cluster.
- **Example:**

-

kubectl get namespaces

Creating and Managing Resources

6. Creating a Namespace

- **Command:**

-

kubectl create namespace mynamespace

- **Purpose:** Creates a new namespace called mynamespace.

7. Setting the Current Context to a Namespace

- **Command:**

-

kubectl config set-context --current --namespace=mynamespace

- **Purpose:** Sets the default namespace to mynamespace for subsequent commands.

8. Creating a Deployment

- **YAML Configuration (my-deployment.yaml):**

yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: mydeployment

namespace: mynamespace

spec:

replicas: 2

selector:

matchLabels:

app: myapp

template:

metadata:

labels:

app: myapp

spec:

containers:

- name: nginx

image: nginx

- **Command to Apply the Deployment:**

-

kubectl apply -f my-deployment.yaml

- **Purpose:** Creates or updates a deployment named mydeployment in the mynamespace namespace.

9. Scaling a Deployment

- **Command:**

-

kubectl scale deployment mydeployment --replicas=3

- **Purpose:** Scales the deployment mydeployment to 3 replicas.

10. Checking the Rollout Status of a Deployment

- **Command:**

-

kubectl rollout status deployment mydeployment

- **Purpose:** Checks the status of the rollout for the deployment mydeployment.

Inspecting and Debugging

5. Describing a Pod

- **Command:**

-

kubectl describe pod mypod

- **Purpose:** Displays detailed information about the pod mypod.

6. Fetching Logs from a Pod

- **Command:**

-

kubectl logs mypod

- **Purpose:** Retrieves and displays logs from the container in the pod mypod.

7. Executing a Command in a Pod

- **Command:**

-

kubectl exec mypod -- /bin/ba-

- **Purpose:** Executes a command within the container of the pod mypod, such as opening a ba- -ell.

8. Listing All Events

- **Command:**

-

kubectl get events

- **Purpose:** Lists all events in the current namespace, useful for debugging.

Advanced Commands

4. -ow Resource Usage of Nodes

- **Command:**

-

kubectl top nodes

- **Purpose:** Displays the current CPU and memory usage of nodes in the cluster.

5. -ow Resource Usage of Pods

- **Command:**

-

kubectl top pods

- **Purpose:** Displays the current CPU and memory usage of pods in the cluster.

6. List All Pods Across All Namespaces

- **Command:**

-

kubectl get pods --all-namespaces

- **Purpose:** Lists all pods across all namespaces for a cluster-wide view.

Example Workflow

8. Create a Namespace

- **Command:**

-

```
kubectrl create namespace mynamespace
```

9. **Switch to the New Namespace**

- **Command:**

-

```
kubectrl config set-context --current --namespace=mynamespace
```

10. **Create a Deployment**

- **Command:**

-

```
kubectrl apply -f my-deployment.yaml
```

11. **Verify the Deployment**

- **Command:**

-

```
kubectrl get deployments
```

12. **Scale the Deployment**

- **Command:**

-

```
kubectrl scale deployment mydeployment --replicas=3
```

13. **Check the Rollout Status**

- **Command:**

-

```
kubectrl rollout status deployment mydeployment
```

14. **Inspect the Pods**

- **Command:**

-

```
kubectrl get pods
```